

Exercise ATags and Sets

Address	Tag	Set
0x0040_17c0	0x00401	496
0x0040_17c4	0x00401	497
0x0040_17c8	0x00401	498
0x0040_17cc	0x00401	499
.	.	.
.	.	.
.	.	.
0x0040_27b8	0x00402	494
0x0040_27bc	0x00402	495
0x0040_27c0	0x00402	496
0x0040_27c4	0x00402	497
0x0040_27c8	0x00402	498
0x0040_27cc	0x00402	499

I-Cache State

Set	Valid	Tag	Instruction
494	1	0x00402	0x1253_0004
495	1	0x00402	0x0240_2021
496	1	0x00401	0x8c99_0000
497	1	0x00402	0x2652_0004
498	1	0x00402	0x1653_fffc
499	1	0x00402	0x2414_0000

Instruction Fetches on the Next Page...

Instruction Fetches

Address	Tag	Set	Action
0x0040_27b8	0x00402	494	I-cache hit - no I-cache update
0x0040_27bc	0x00402	495	I-cache miss - instruction 0x0240_2021 is copied into instruction field in set 495, V-bit in that set is changed to 1, tag to 0x00402
0x0040_27c0	0x00402	496	I-cache miss - instruction 0x0c10_0f50 is copied into instruction field in set 496, V-bit in that set is changed to 1, tag to 0x00402
0x0040_17c0	0x00401	496	I-cache miss - instruction 0x8c99_0000 is copied into instruction field in set 496, V-bit in that set is changed to 1, tag to 0x00401
0x0040_17c4	0x00401	497	I-cache hit - no I-cache update
0x0040_17c8	0x00401	498	I-cache hit - no I-cache update
0x0040_17cc	0x00401	499	I-cache hit - no I-cache update
0x0040_27c4	0x00402	497	I-cache miss - instruction 0x2652_0004 is copied into instruction field in set 497, V-bit in that set is changed to 1, tag to 0x00402
0x0040_27c8	0x00402	498	I-cache miss - instruction 0x1653_ffff is copied into instruction field in set 498, V-bit in that set is changed to 1, tag to 0x00402
0x0040_27bc	0x00402	495	I-cache hit - no I-cache update
0x0040_27c0	0x00402	496	I-cache miss - instruction 0x0c10_0f50 is copied into instruction field in set 496, V-bit in that set is changed to 1, tag to 0x00402
0x0040_17c0	0x00401	496	I-cache miss - instruction 0x8c99_0000 is copied into instruction field in set 496, V-bit in that set is changed to 1, tag to 0x00401
0x0040_17c4	0x00401	497	I-cache miss - instruction 0x0019_c023 is copied into instruction field in set 497, V-bit in that set is changed to 1, tag to 0x00401
0x0040_17c8	0x00401	498	I-cache miss - instruction 0xac98_0000 is copied into instruction field in set 498, V-bit in that set is changed to 1, tag to 0x00401
0x0040_17cc	0x00401	499	I-cache hit - no I-cache update
0x0040_27c4	0x00402	497	I-cache miss - instruction 0x2652_0004 is copied into instruction field in set 497, V-bit in that set is changed to 1, tag to 0x00402
0x0040_27c8	0x00402	498	I-cache miss - instruction 0x1653_ffff is copied into instruction field in set 498, V-bit in that set is changed to 1, tag to 0x00402
0x0040_27cc	0x00402	499	I-cache miss - instruction 0x2414_0000 is copied into instruction field in set 499, V-bit in that set is changed to 1, tag to 0x00402

Exercise B

3. $C = 32 \text{ KB} = 32000 \text{ bytes} = 2^{15} \text{ bytes}$
 words per block = 8 words/block = 2^3 words/block
 bytes per word = 4 bytes/word = 2^2 bytes/word
 $S = ?$

a) Number of Sets:

$$C = S \times \text{words per block} \times \text{bytes per word}$$

$$2^{15} = S \times 2^3 \frac{\text{words}}{\text{block}} \times 2^2 \frac{\text{bytes}}{\text{word}}$$

$$S = \frac{2^{15} \text{ bytes}}{2^3 \frac{\text{words}}{\text{block}} \times 2^2 \frac{\text{bytes}}{\text{word}}} = 2^{15-3-2} = 1024 \text{ blocks}$$

b) Address Split:

The address split is:

- byte offset, $\log_2 4 = 2 \text{ bits}$;
- block offset, $\log_2 8 = 3 \text{ bits}$;
- set bits, $\log_2 1024 = 10 \text{ bits}$;
- tag bits, $32 - 10 - 3 - 2 = 17 \text{ bits}$.

c) Bits Per Block:

A block has:

- 1 V-bit,
- 17 tag bits,
- 8 words/block \times 32 bits/word = 256 data bits.

Thus, there are $1 + 17 + 256 = 274$ bits per block.

SRAM Cells:

$$\text{Total number of bits} = 274 \frac{\text{bits}}{\text{block}} \times 1024 \text{ blocks} = 280576 \text{ bits}$$

Thus, 280576 SRAM cells are needed for the whole cache.

4. bytes per block = 64 bytes/block
 words per block = ?
 bytes per word = 8 bytes/word = 2^3 bytes/word
 S = ?

a) $C = 64 \text{ KB} = 64000 \text{ bytes} = 2^{16} \text{ bytes}$

Words Per Block:

$$\text{Bytes per block} = \text{bytes per word} \times \text{words per block}$$

$$64 \frac{\text{bytes}}{\text{block}} = 8 \frac{\text{bytes}}{\text{word}} \times \text{words per block}$$

$$\text{Words per block} = \frac{64 \frac{\text{bytes}}{\text{block}}}{8 \frac{\text{bytes}}{\text{word}}} = 8 \frac{\text{words}}{\text{block}} = 2^3 \frac{\text{words}}{\text{block}}$$

Number of Sets:

$$C = S \times \text{words per block} \times \text{bytes per word}$$

$$2^{16} = S \times 2^3 \frac{\text{words}}{\text{block}} \times 2^3 \frac{\text{bytes}}{\text{word}}$$

$$S = \frac{2^{16} \text{ bytes}}{2^3 \frac{\text{words}}{\text{block}} \times 2^3 \frac{\text{bytes}}{\text{word}}} = 2^{16-3-3} = 1024 \text{ blocks}$$

Address Split:

The address split is:

- byte offset, $\log_2 8 = 3 \text{ bits}$;
- block offset, $\log_2 8 = 3 \text{ bits}$;
- set bits, $\log_2 1024 = 10 \text{ bits}$;
- tag bits, $44 - 10 - 3 - 3 = 28 \text{ bits}$.

43	16	15	6	5	3	2	0
Search Tag		Set Bit		Block Offset		Byte Offset	
28 bits		10 bits		3 bits		3 bits	

b) $C = 2 \text{ MB} = 2000000 \text{ bytes} = 2^{21} \text{ bytes}$

Words Per Block:

$$\text{Bytes per block} = \text{bytes per word} \times \text{words per block}$$

$$64 \frac{\text{bytes}}{\text{block}} = 8 \frac{\text{bytes}}{\text{word}} \times \text{words per block}$$

$$\text{Words per block} = \frac{64 \frac{\text{bytes}}{\text{block}}}{8 \frac{\text{bytes}}{\text{word}}} = 8 \frac{\text{words}}{\text{block}} = 2^3 \frac{\text{words}}{\text{block}}$$

Number of Sets:

$$C = S \times \text{words per block} \times \text{bytes per word}$$

$$2^{21} = S \times 2^3 \frac{\text{words}}{\text{block}} \times 2^3 \frac{\text{bytes}}{\text{word}}$$

$$S = \frac{2^{21} \text{ bytes}}{2^3 \frac{\text{words}}{\text{block}} \times 2^3 \frac{\text{bytes}}{\text{word}}} = 2^{21-3-3} = 32768 \text{ blocks}$$

Address Split:

The address split is:

- byte offset, $\log_2 8 = 3 \text{ bits}$;
- block offset, $\log_2 8 = 3 \text{ bits}$;
- set bits, $\log_2 32768 = 15 \text{ bits}$;
- tag bits, $44 - 15 - 3 - 3 = 23 \text{ bits}$.

43	21	20	6	5	3	2	0
Search Tag	Set Bit	Block Offset	Byte Offset				
23 bits	15 bits	3 bits	3 bits				

c) Bits Per Block:

A block has: 1 V-bit,
 23 tag bits,
 8 words/block \times 64 bits/word = 512 data bits.

Thus, there are $1 + 23 + 512 = 536$ bits per block.

SRAM Cells:

$$\text{Total number of bits} = 536 \frac{\text{bits}}{\text{block}} \times 32768 \text{ blocks} = 17563648 \text{ bits}$$

Thus, 17563648 SRAM cells are needed for the whole cache.

Exercise C

Part II

Capacity:

$C = ?$

words per block = 4 words/block

bytes per word = 4 bytes/word

$S = 256$ blocks

$$C = S \times \text{words per block} \times \text{bytes per word}$$

$$C = 256 \text{ blocks} \times 4 \frac{\text{words}}{\text{block}} \times 4 \frac{\text{bytes}}{\text{word}} = 4096 \text{ bytes}$$

Address Split:

The address split is: byte offset, $\log_2 4 = 2 \text{ bits}$;
 block offset, $\log_2 4 = 2 \text{ bits}$;
 set bits, $\log_2 256 = 8 \text{ bits}$;
 tag bits, $32 - 8 - 2 - 2 = 20 \text{ bits}$.

31	12	11	4	3	2	1	0	
Search Tag				Set Bit		Block Offset		Byte Offset
20 bits				8 bits		2 bits		2 bits

Program Runs:

	heapsort_trace.txt	mergesort_trace.txt
sim1.c	64705 reads 49791 read hits 60419 writes 60401 write hits overall miss rate: 11.9%	104298 reads 88429 read hits 73410 writes 64076 write hits overall miss rate: 14.2%
sim2.c	64705 reads 57046 read hits 60419 writes 60387 write hits overall miss rate: 6.1%	104298 reads 99955 read hits 73410 writes 70863 write hits overall miss rate: 3.9%

There is significant spatial locality of reference in the memory accesses done by both algorithms because, even though both caches are the same size, the 4-word block performs a lot better than the 1-word block cache.

Part III

Capacity:

$C = 4 \text{ KB} = 4096 \text{ bytes} = 2^{12} \text{ bytes}$
 words per block = 16 words/block = 2^4 words/block
 bytes per word = 4 bytes/word = 2^2 bytes/word
 $S = ?$

$$C = S \times \text{words per block} \times \text{bytes per word}$$

$$2^{12} \text{ bytes} = S \times 2^4 \frac{\text{words}}{\text{block}} \times 2^2 \frac{\text{bytes}}{\text{word}}$$

$$S = \frac{2^{12} \text{ bytes}}{2^4 \frac{\text{words}}{\text{block}} \times 2^2 \frac{\text{bytes}}{\text{word}}} = 2^{12-4-2} = 64 \text{ blocks}$$

Address Split:

The address split is:

- byte offset, $\log_2 4 = 2 \text{ bits}$;
- block offset, $\log_2 16 = 4 \text{ bits}$;
- set bits, $\log_2 64 = 6 \text{ bits}$;
- tag bits, $32 - 6 - 4 - 2 = 20 \text{ bits}$.

31	11	6 5	2 1	0
Search Tag	Set Bit	Block Offset	Byte Offset	
20 bits	6 bits	4 bits	2 bits	

Program Runs:

	heapsort_trace.txt	mergesort_trace.txt
sim3.c	64705 reads 59875 read hits 60419 writes 60324 write hits overall miss rate: 3.9%	104298 reads 101949 read hits 73410 writes 72243 write hits overall miss rate: 2.0%

Part II Source File:

```
// sim2.c
// ENCM 369 Winter 2018 Lab 10 Exercise C
// Author: S. Norman
// Edited by: R. Lim
//
// If you build an executable using gcc -Wall sim2.c -o sim2
// you can run it by redirecting input to come from a data file,
// as in
//      ./sim2 < heapsort_trace.txt

#include <stdio.h>
#include <stdlib.h>

int read_one_line(unsigned *p)
    // Read one line of the input stream.
    // Return value is normally 'r' or 'w' to indicate read or write.
    // In that case, *p contains the address read from the input line.
    // Return value is 'e' to indicate that input failed at the end of
    // the input stream.

{
    int nscan, rw;
    char buf[2];

    nscan = scanf("%1s%x", buf, p);
    if (nscan == EOF)
        return 'e';                                /* indicate end-of-file */
    else if (nscan != 2) {
        fprintf(stderr, "Format error in input stream.\n");
        exit(1);
    }

    rw = buf[0];
    if (rw != 'r' && rw != 'w') {
        fprintf(stderr, "Read/write character was neither r nor w.\n");
        exit(1);
    }
    return rw;
}

// These two arrays keep track of all the V-bits and stored tags in
// the array. We don't need an array for data to count hits and misses.
// Because these arrays are external variables, it's safe to assume
// it will be initialized to all zeros before main starts.
char v_array[256];
unsigned tag_array[256];

int main(void)
```

```
{
    int read_count = 0, read_hits = 0;
    int write_count = 0, write_hits = 0;
    int access_count, miss_count;
    int rw;
    unsigned address, tag, set;
    int hit;

    while (1) {
        rw = read_one_line(&address);
        if (rw == 'e') break;

        set = (address & 0xff0) >> 4; // bits 11-4
        tag = address >> 12;          // bits 31-12

        // Note: Next line results in either hit == 1 or hit == 0.
        hit = v_array[set] == 1 && tag_array[set] == tag;
        if (rw == 'r') {
            read_count++;
            read_hits += hit;
        }
        else {
            write_count++;
            write_hits += hit;
        }
        if (!hit) { // On a miss, update V-bit and tag.
            v_array[set] = 1;
            tag_array[set] = tag;
        }
    }

    printf("%d reads\n", read_count);
    printf("%d read hits\n", read_hits);
    printf("%d writes\n", write_count);
    printf("%d write hits\n", write_hits);

    access_count = read_count + write_count;
    miss_count = access_count - read_hits - write_hits;
    printf("overall miss rate: %.1f%%\n",
        100.0 * (double) miss_count / access_count);

    return 0;
}
```

Part III Source File:

```
// sim3.c
// ENCM 369 Winter 2018 Lab 10 Exercise C
// Author: S. Norman
// Edited by: R. Lim
//
// If you build an executable using gcc -Wall sim3.c -o sim3
// you can run it by redirecting input to come from a data file,
// as in
//      ./sim3 < heapsort_trace.txt

#include <stdio.h>
#include <stdlib.h>

int read_one_line(unsigned *p)
    // Read one line of the input stream.
    // Return value is normally 'r' or 'w' to indicate read or write.
    // In that case, *p contains the address read from the input line.
    // Return value is 'e' to indicate that input failed at the end of
    // the input stream.

{
    int nscan, rw;
    char buf[2];

    nscan = scanf("%1s%x", buf, p);
    if (nscan == EOF)
        return 'e';                                /* indicate end-of-file */
    else if (nscan != 2) {
        fprintf(stderr, "Format error in input stream.\n");
        exit(1);
    }

    rw = buf[0];
    if (rw != 'r' && rw != 'w') {
        fprintf(stderr, "Read/write character was neither r nor w.\n");
        exit(1);
    }
    return rw;
}

// These two arrays keep track of all the V-bits and stored tags in
// the array. We don't need an array for data to count hits and misses.
// Because these arrays are external variables, it's safe to assume
// it will be initialized to all zeros before main starts.
char v_array[64];
unsigned tag_array[64];

int main(void)
```

```

{
    int read_count = 0, read_hits = 0;
    int write_count = 0, write_hits = 0;
    int access_count, miss_count;
    int rw;
    unsigned address, tag, set;
    int hit;

    while (1) {
        rw = read_one_line(&address);
        if (rw == 'e') break;

        set = (address & 0xfc0) >> 6; // bits 11-6
        tag = address >> 12;          // bits 31-12

        // Note: Next line results in either hit == 1 or hit == 0.
        hit = v_array[set] == 1 && tag_array[set] == tag;
        if (rw == 'r') {
            read_count++;
            read_hits += hit;
        }
        else {
            write_count++;
            write_hits += hit;
        }
        if (!hit) { // On a miss, update V-bit and tag.
            v_array[set] = 1;
            tag_array[set] = tag;
        }
    }

    printf("%d reads\n", read_count);
    printf("%d read hits\n", read_hits);
    printf("%d writes\n", write_count);
    printf("%d write hits\n", write_hits);

    access_count = read_count + write_count;
    miss_count = access_count - read_hits - write_hits;
    printf("overall miss rate: %.1f%%\n",
        100.0 * (double) miss_count / access_count);

    return 0;
}

```